

## Lab9: PCA + Clustering

```
In [67]: ID = 1631938
        Name = 'MIN SOE HTUT'
```

### Task 1: Pre-done: Loading the data and preparing the data

```
In [68]: from datetime import datetime
import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
df = pd.read_csv('https://raw.githubusercontent.com/bpfa/data_for_compx310_2024/main/marketing_campaign.csv', sep=',')
# Display the first few rows to inspect the data
print(df.columns)

Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome',
       'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits',
       'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
       'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
       'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth',
       'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
       'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'],
      dtype='object')
```

### Creating new features for customer profiles

```
In [69]: # Creates a new field to store the age of the customer
df['Age']=2022-df['Year_Birth']

# Recodes the customer's education level to numeric form (0: high-school, 1:
diploma, 2: bachelors, 3: masters, and 4: doctorates)
df["Education"].replace({"Basic":0,"2n Cycle":1, "Graduation":2, "Master":3,
"PhD":4},inplace=True)

# Recodes the customer's marital status to numeric form (0: not living with a
partner, 1: living with a partner)
df['Marital_Status'].replace({"Married":1, "Together":1, "Absurd":0, "Widow":
0, "YOLO":0, "Divorced":0, "Single":0,"Alone":0},inplace=True)

# creates a new field to store the number of children in the household
df['Children']=df['Kidhome']+df['Teenhome']

# creates a new field to store the household size
df['Family_Size']=df['Marital_Status']+df['Children']+1
```

### Creating new features for customer spending behavior

```
In [70]: # creates a new field to store the total spending of the customer
df['Total_Spending']=df["MntWines"]+ df["MntFruits"]+ df["MntMeatProducts"]+ d
f["MntFishProducts"]+ df["MntSweetProducts"]+ df["MntGoldProds"]

# creates subsequent fields to store the spending proportion for each product
by the customer
df['Prop_Wines']=df["MntWines"]/df["Total_Spending"]
df['Prop_Fruits']=df["MntFruits"]/df["Total_Spending"]
df['Prop_MeatProducts']=df["MntMeatProducts"]/df["Total_Spending"]
df['Prop_FishProducts']=df["MntFishProducts"]/df["Total_Spending"]
df['Prop_SweetProducts']=df["MntSweetProducts"]/df["Total_Spending"]
df['Prop_GoldProds']=df["MntGoldProds"]/df["Total_Spending"]
```

### Additional customer features

```
In [71]: # Converts the customer start date into a useful feature representing how long
they have been a customer
df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'], dayfirst=True)
# Get the current date
today = datetime.today()
# Calculate how long each customer has been with the company
df['Days_as_Customer'] = (today - df['Dt_Customer']).dt.days
# Combine the number of offers a customer responded to into a single feature
df['Offers_Responded_To'] = df['AcceptedCmp1'] + df['AcceptedCmp2'] + df['Acce
ptedCmp3'] + df['AcceptedCmp4'] + df['AcceptedCmp5'] + df['Response']
```

### Cleaning the data by removing outliers

```
In [72]: # Remove outliers when we do customer segmentation, as we are more interested
in the general population rather than the outliers
df = df[(df["Age"]<90)]
df = df[(df["Income"]<110000)]
df = df[(df["NumWebVisitsMonth"]<11)]
df = df[(df["NumWebPurchases"]<20)]
df = df[(df["NumCatalogPurchases"]<20)]
```

### Dropping irrelevant or unhelpful fields for clustering

```
In [73]: fields_to_drop=['ID','Year_Birth','Dt_Customer','Z_CostContact','Z_Revenue','A
cceptedCmp1',
                        'AcceptedCmp2','AcceptedCmp3','AcceptedCmp4','AcceptedCmp5','Respon
se','Complain',
                        'MntFruits','MntWines','MntMeatProducts','MntFishProducts','MntSweetP
roducts','MntGoldProds']
df.drop(fields_to_drop,axis=1,inplace=True)

# Remove any remaining missing values
df.dropna(inplace=True)
```

### 3D scatter plots

```
In [74]: def scatter_3d(x,y,z,c=None):
fig = plt.figure(figsize=(10,8))
ax = plt.subplot(111, projection='3d', label="bla")
ax.scatter(x, y, z, s=40, c=c, marker='o',cmap=plt.cm.viridis)
ax.set_title("The Plot Of The Clusters")
plt.show()
```

### Task 2: PCA

a) Apply StandardScaler preprocessing on the dataframe df, and assigned the fit\_transformed values as df\_scaled

```
In [75]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)
```


b) Apply PCA transformation on df\_scaled using all 23 components

```
In [76]: from sklearn.decomposition import PCA
pca = PCA(n_components=23)
df_pca = pca.fit_transform(df_scaled)
```

### c) Plot the cumulative

```
In [77]: cumsum = np.cumsum(pca.explained_variance_ratio_)

plt.figure(figsize=(10, 8))
plt.plot(range(1, len(cumsum) + 1), cumsum, marker='o')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Variance Explained')
plt.title('Cumulative Variance Explained by Principal Components')
plt.grid(True)
plt.show()
```



```
In [78]: # Determine the number of components needed for 80% variance
n_components = np.argmax(cumsum >= 0.80) + 1
print(f'Number of components explaining 80% of the variance: {n_components}')
```

Number of components explaining 80% of the variance: 10

According to the plot 10 Principal Components is needed for a least 80% of the variance is explained since for 82% , 10 Principal Components is needed.

**d) Redo the PCA transformation on `df_scaled` using the same number of components as the value from 2c**

```
In [79]: pca = PCA(n_components=n_components)
df_transformed = pca.fit_transform(df_scaled)
```

**e) Visualize the first 3 components of `df_transformed` using a 3d scatter plot**


```
In [80]: def scatter_3d(x, y, z, c=None):
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
sc = ax.scatter(x, y, z, s=40, c=c, marker='o', cmap=plt.cm.viridis)
if c is not None:
fig.colorbar(sc)
ax.set_title("3D Scatter Plot of PCA Components")
plt.show()

# Visualize the first 3 PCA components
scatter_3d(df_transformed[:, 0], df_transformed[:, 1], df_transformed[:, 2])
```

<ipython-input-80-7f4b135a8134>:4: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored

```
sc = ax.scatter(x, y, z, s=40, c=c, marker='o', cmap=plt.cm.viridis)
```

3D Scatter Plot of PCA Components



### Task 3: KMeans Clustering

a) Apply KMeans clustering to `df_transformed` and measure the inertia score for `n_clusters` between 2 and 10

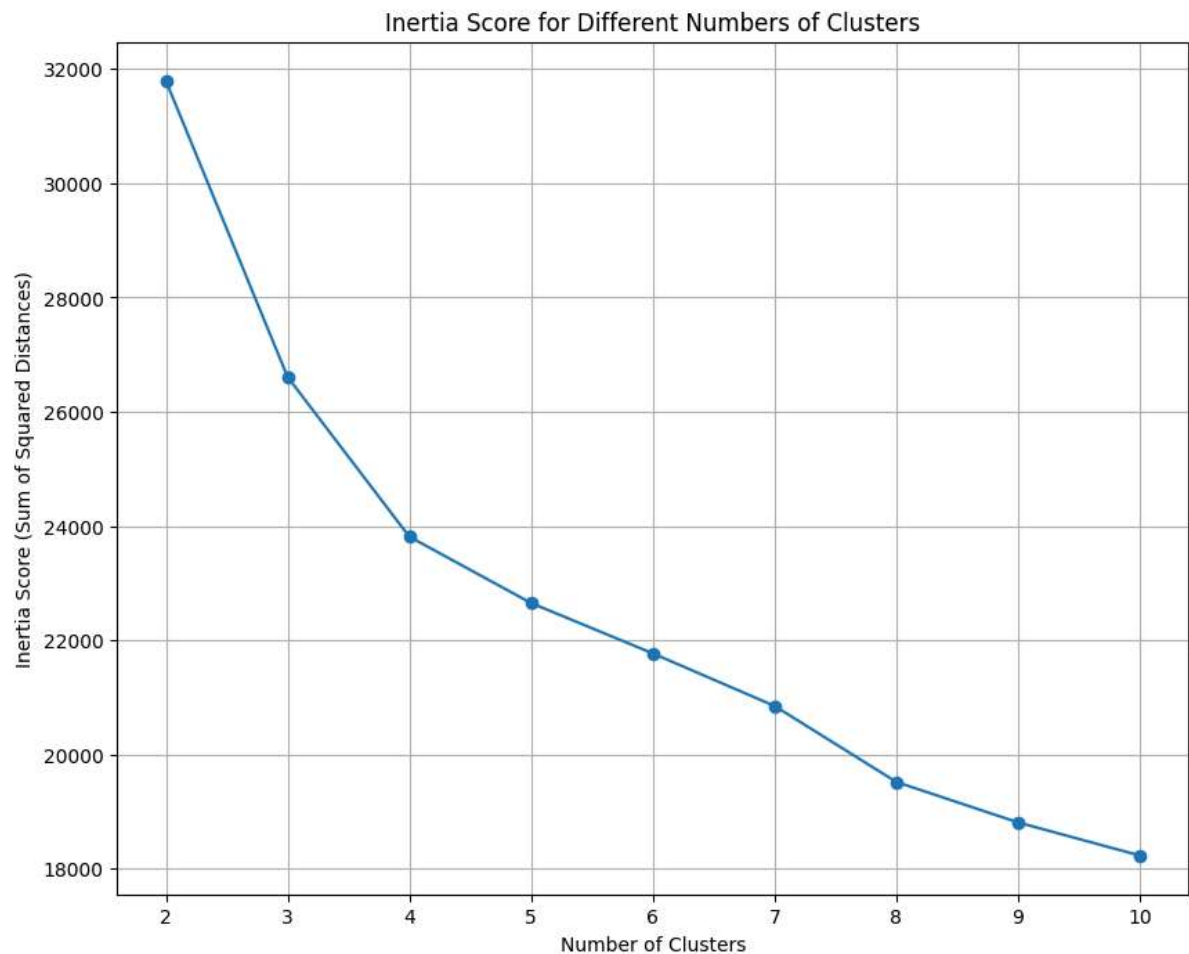
```
In [81]: from sklearn.cluster import KMeans

inertia_scores = []
cluster_range = range(2, 11)

for i in cluster_range:
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(df_transformed)
    inertia_scores.append(kmeans.inertia_)
```

b) Plot the inertia score.

```
In [82]: plt.figure(figsize=(10, 8))
plt.plot(cluster_range, inertia_scores, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia Score (Sum of Squared Distances)')
plt.title('Inertia Score for Different Numbers of Clusters')
plt.grid(True)
plt.show()
```



**What would you say is the best number of clusters for this dataset? Why?**

The best number of the clusters is 4 because according to the plot, it can be seen that the number of clusters being 4 and the inertia score being just below 24000 is the elbow point where both the number of clusters and inertia score are balanced and equally lowest. Before this point, the number of clusters is low, but the inertia score is pretty high and then the inertia score decreases sharply but the number of clusters increases. Therefore, this elbow point is the best point for both the number of clusters and inertia score.

**c) Find the cluster labels of KMeans(n\_clusters=4) on df\_transformed, (It doesn't matter what your answer for 3b is, set n\_clusters=4)**

```
In [83]: kmeans = KMeans(n_clusters=4, random_state= ID)
cluster_labels = kmeans.fit_predict(df_transformed)
```


**d) Assign the cluster labels as df['Clusters']**

```
In [84]: df['Clusters'] = cluster_labels
```

**e) Visualize the first 3 components of df\_transformed using a 3d scatter plot, with the data points coloured according to the clusters**




```
In [85]: scatter_3d(df_transformed[:, 0], df_transformed[:, 1], df_transformed[:, 2], c  
              =df['Clusters'])
```



#### Task 4: Interpretating our results

a) Do a scatter plot between `df['Income']` and `df['Total_Spending']` and colour the data points according the clusters.

```
In [86]: pl = sns.scatterplot(data=df, x="Total_Spending", y="Income", hue="Clusters")  
pl.set_title("Cluster's Profile Based On Income And Spending")  
plt.show()
```



b) Do a boxplot by clusters, for each of these following fields describing the attributes of the customer ['Age', 'Education', 'Marital\_Status', 'Income', 'Children', 'Family\_Size', 'Kidhome', 'Teenhome']

```

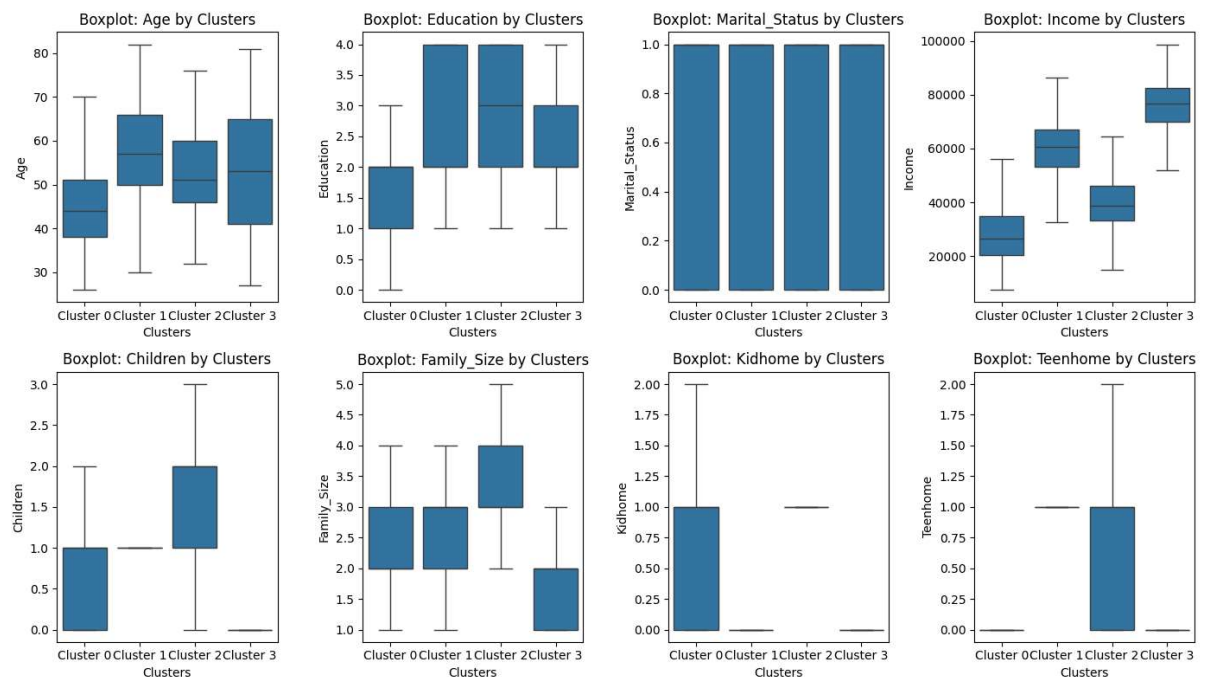
In [87]: import seaborn as sns
import matplotlib.pyplot as plt

# Get the number of unique clusters
n_clusters = df['Clusters'].nunique()

# Boxplots for customer attributes by clusters
attributes = ['Age', 'Education', 'Marital_Status', 'Income', 'Children', 'Family_Size', 'Kidhome', 'Teenhome']

plt.figure(figsize=(14, 8))
for i, attribute in enumerate(attributes):
    plt.subplot(2, 4, i + 1) # Create subplots in a 2x4 grid
    sns.boxplot(x='Clusters', y=attribute, data=df, showfliers=False)
    plt.title(f'Boxplot: {attribute} by Clusters')
    plt.xticks(range(n_clusters), [f'Cluster {i}' for i in range(n_clusters)])
plt.tight_layout()
plt.show()

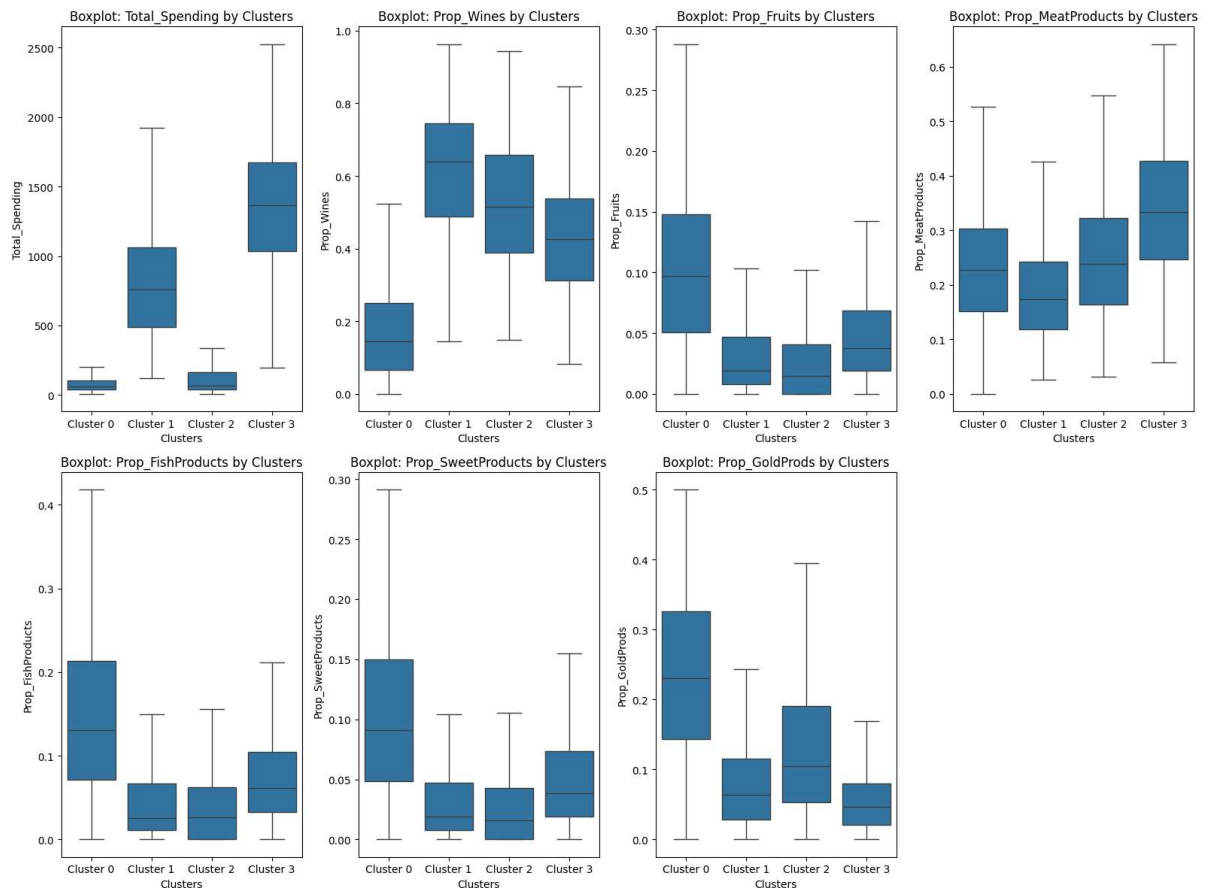
```



c) Do a boxplot by clusters, for each of these following fields describing the customer preference ['Total\_Spending','Prop\_Wines','Prop\_Fruits','Prop\_MeatProducts','Prop\_FishProducts','Prop\_SweetProdu

```
In [88]: # Boxplots for customer preferences by clusters
preferences = ['Total_Spending', 'Prop_Wines', 'Prop_Fruits', 'Prop_MeatProducts',
               'Prop_FishProducts', 'Prop_SweetProducts', 'Prop_GoldProds']

plt.figure(figsize=(16, 12))
for i, preference in enumerate(preferences):
    plt.subplot(2, 4, i + 1)
    sns.boxplot(x='Clusters', y=preference, data=df, showfliers=False)
    plt.title(f'Boxplot: {preference} by Clusters')
    plt.xticks(range(n_clusters), [f'Cluster {i}' for i in range(n_clusters)])
plt.tight_layout()
plt.show()
```

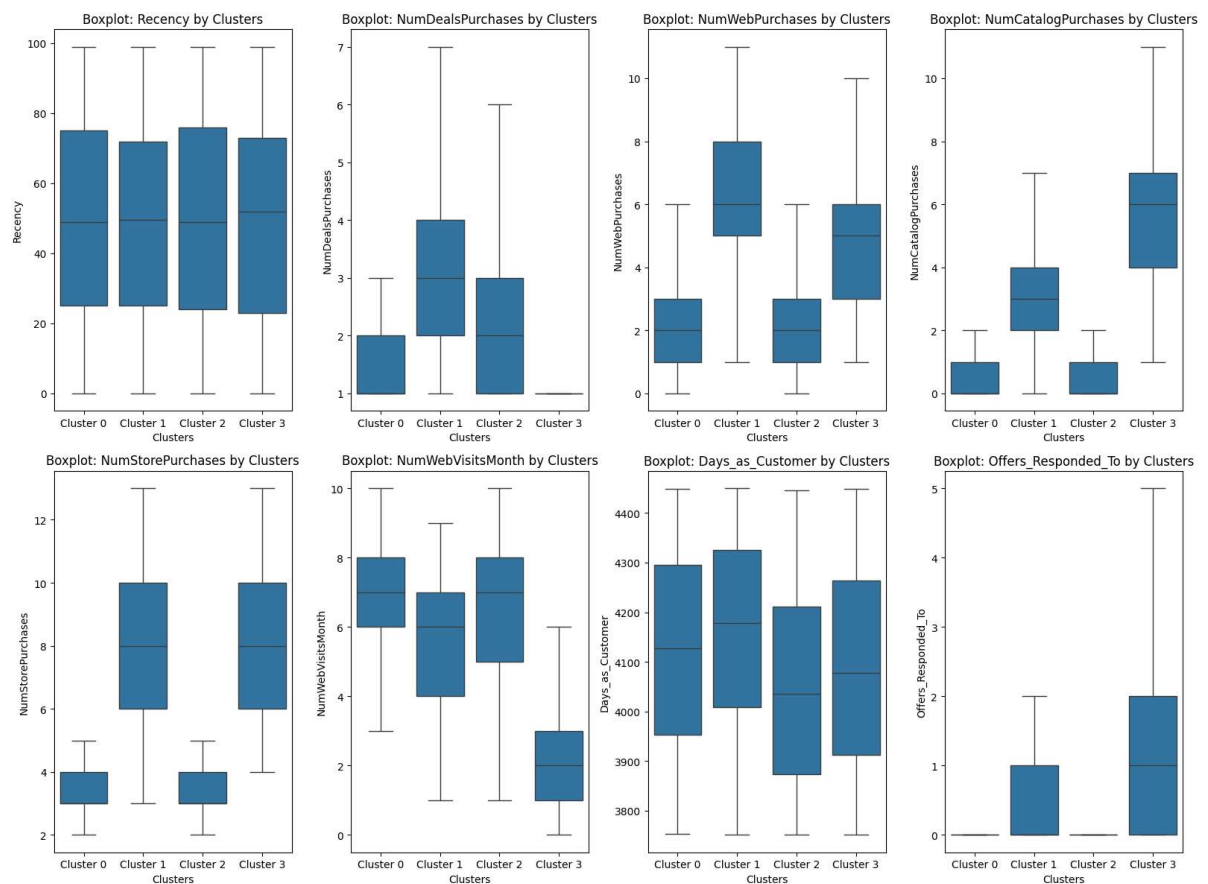


d) Do a boxplot by clusters, for each of these following fields describing the customer behaviour ['Recency', 'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth', 'Days\_as\_Customer', 'Offers\_Responded\_To']

```
In [89]: # Get the number of unique clusters dynamically
n_clusters = df['Clusters'].nunique()

# Boxplots for customer behavior by clusters
behaviors = ['Recency', 'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
             'NumStorePurchases', 'NumWebVisitsMonth', 'Days_as_Customer', 'Offers_Responded_To']

plt.figure(figsize=(16, 12))
for i, behavior in enumerate(behaviors):
    plt.subplot(2, 4, i + 1)
    sns.boxplot(x='Clusters', y=behavior, data=df, showfliers=False)
    plt.title(f'Boxplot: {behavior} by Clusters')
    plt.xticks(range(n_clusters), [f'Cluster {i}' for i in range(n_clusters)])
# Handle different numbers of clusters
plt.tight_layout()
plt.show()
```



e) Obtain the means grouped by the cluster, for each of the fields. (See hints for a 1-liner code)

```
In [90]: # Calculate the means of the features grouped by Clusters
cluster_means = df.groupby('Clusters').mean().T
print(cluster_means)
```

Clusters	0	1	2	3
Education	1.664319	2.656151	2.753647	2.506796
Marital_Status	0.624413	0.675079	0.675851	0.592233
Income	28299.960094	60160.159306	39556.867099	75822.452427
Kidhome	0.678404	0.156151	0.925446	0.023301
Teenhome	0.096244	0.968454	0.696921	0.048544
Recency	49.117371	48.676656	49.264182	49.225243
NumDealsPurchases	1.821596	3.488959	2.470016	1.046602
NumWebPurchases	2.173709	6.361199	2.406807	4.893204
NumCatalogPurchases	0.591549	3.361199	0.619125	5.899029
NumStorePurchases	3.293427	7.973186	3.411669	8.300971
NumWebVisitsMonth	6.546948	5.550473	6.385737	2.642718
Age	45.730047	57.913249	53.076175	53.264078
Children	0.774648	1.124606	1.622366	0.071845
Family_Size	2.399061	2.799685	3.298217	1.664078
Total_Spending	111.453052	807.957413	117.638574	1363.409709
Prop_Wines	0.167191	0.616963	0.524853	0.433415
Prop_Fruits	0.106278	0.034664	0.025610	0.049006
Prop_MeatProducts	0.233265	0.185796	0.251463	0.334373
Prop_FishProducts	0.151586	0.046297	0.041214	0.074083
Prop_SweetProducts	0.106718	0.035482	0.026846	0.050715
Prop_GoldProds	0.234962	0.080799	0.130014	0.058409
Days_as_Customer	4123.093897	4154.750789	4054.682334	4090.203883
Offers_Responded_To	0.180751	0.386435	0.183144	1.081553